



Machine Learning

Linear Regression with
multiple variables

Multiple features

Multiple features (variables).

Size (feet ²)	Price (\$1000)
 x	y 
2104	460
1416	232
1534	315
852	178
...	...

$$\underline{h_{\theta}(x) = \theta_0 + \theta_1 x}$$

Multiple features (variables).

→ Size (feet²) Number of bedrooms Number of floors Age of home (years) Price (\$1000)

x_1	x_2	x_3	x_4	y
2104	5	1	45	460
→ 1416	3	2	40	232
1534	3	2	30	315
852	2	1	36	178
...

Notation:

- n = number of features $n = 4$
- $x^{(i)}$ = input (features) of i^{th} training example.
- $x_j^{(i)}$ = value of feature j in i^{th} training example.

$m = 47$

$x^{(2)} = \begin{bmatrix} 1416 \\ 3 \\ 2 \\ 40 \end{bmatrix}$

$x_3^{(2)} = 2$

Hypothesis:

Previously:

~~$$h_{\theta}(x) = \theta_0 + \theta_1 x$$~~

$$k_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4$$

E.g. $\underline{h_{\theta}(x)} = \underline{80} + \underline{0.1x_1} + \underline{0.01x_2} + \underline{3x_3} - \underline{2x_4}$
↑ ↑ ↑
age

$$\rightarrow h_{\theta}(x) = \underline{\theta_0} + \underline{\theta_1}x_1 + \underline{\theta_2}x_2 + \dots + \underline{\theta_n}x_n$$

For convenience of notation, define $\boxed{x_0 = 1}$ ($x_0^{(i)} = 1$)

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$h_{\theta}(x) = \underset{\substack{\uparrow \\ = 1}}{\theta_0}x_0 + \theta_1x_1 + \dots + \theta_nx_n$$


$$= \boxed{\theta^T x}$$

$\underbrace{\begin{bmatrix} \theta_0 & \theta_1 & \dots & \theta_n \end{bmatrix}}_{\theta^T}$
 $(n+1) \times 1$
 matrix

\swarrow

$\begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix}$
 x

Multivariate linear regression. \leftarrow



Machine Learning

Linear Regression with
multiple variables

Gradient descent for
multiple variables

Hypothesis: $h_{\theta}(x) = \theta^T x = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$ $\rightarrow x_0 = 1$

Parameters: $\theta_0, \theta_1, \dots, \theta_n$ θ $n+1$ -dimensional vector

Cost function:

$$~~J(\theta_0, \theta_1, \dots, \theta_n)~~ = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$J(\theta)$

Gradient descent:

Repeat {

$\rightarrow \theta_j := \theta_j - \alpha \left[\frac{\partial}{\partial \theta_j} J(\theta_0, \dots, \theta_n) \right] J(\theta)$
(simultaneously update for every $j = 0, \dots, n$)

Gradient Descent

Previously ($n=1$):

Repeat {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\frac{\partial}{\partial \theta_0} J(\theta)$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)}$$

(simultaneously update θ_0, θ_1)

}

New algorithm ($n \geq 1$):

Repeat {

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

$$\frac{\partial}{\partial \theta_j} J(\theta)$$

(simultaneously update θ_j for $j = 0, \dots, n$)


}

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_1^{(i)}$$

$$\theta_2 := \theta_2 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_2^{(i)}$$

...



Machine Learning

Linear Regression with multiple variables


Gradient descent in practice I: Feature Scaling

Feature Scaling

Idea: Make sure features are on a similar scale.

E.g. x_1 = size (0-2000 feet²)


x_2 = number of bedrooms (1-5)



$$\rightarrow x_1 = \frac{\text{size (feet}^2\text{)}}{2000}$$

$$\rightarrow x_2 = \frac{\text{number of bedrooms}}{5}$$

$$0 \leq x_1 \leq 1 \quad 0 \leq x_2 \leq 1$$



Feature Scaling

Get every feature into approximately a $[-1 \leq x_i \leq 1]$ range.

$$x_0 = 1$$

$$0 \leq x_1 \leq 3 \quad \checkmark$$

$$-2 \leq x_2 \leq 0.5 \quad \checkmark$$

$$-100 \leq x_3 \leq \boxed{100} \quad \times$$

$$-0.0001 \leq x_4 \leq \boxed{0.0001} \quad \times$$

$$\boxed{-1 \leq x_i \leq 1} \text{ range.}$$

$$-3 \text{ to } 3 \quad \checkmark$$

$$-\frac{1}{3} \text{ to } \frac{1}{3} \quad \checkmark$$

Mean normalization

Replace x_i with $x_i - \mu_i$ to make features have approximately zero mean
(Do not apply to $x_0 = 1$).

E.g. $x_1 = \frac{\text{size} - 1000}{2000}$

Average size ≈ 1000

$$x_2 = \frac{\# \text{bedrooms} - 2}{5}$$

1-5 bedrooms

$$-0.5 \leq x_1 \leq 0.5$$


$$-0.5 \leq x_2 \leq 0.5$$

$$x_1 \leftarrow$$

$$\frac{x_1 - \mu_1}{s_1}$$

μ_1 ← avg value of x_1 in training set
 s_1 ← range (max - min) (or standard deviation)

$$x_2 \leftarrow \frac{x_2 - \mu_2}{s_2}$$



Machine Learning

Linear Regression with multiple variables


Gradient descent in practice II: Learning rate

Gradient descent

→ $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$

- “Debugging”: How to make sure gradient descent is working correctly.
- How to choose learning rate α .

Making sure gradient descent is working correctly.




→ Example automatic convergence test:

→ Declare convergence if $J(\theta)$ decreases by less than 10^{-3} in one iteration.

ϵ


Making sure gradient descent is working correctly.



- For sufficiently small α , $J(\theta)$ should decrease on every iteration. ←
- But if α is too small, gradient descent can be slow to converge.

Summary:

- If α is too small: slow convergence.
- If α is too large: $J(\theta)$ may not decrease on every iteration; may not converge. (Slow converge also possible)



To choose α , try


$\dots, \underline{0.001}, \underline{0.003}, \underline{0.01}, \underline{0.03}, \underline{0.1}, \underline{0.3}, \underline{1}, \dots$

Handwritten annotations below the values:

- Under 0.001: \nearrow (curved arrow pointing up and right)
- Under 0.003: \nearrow (curved arrow pointing up and right)
- Under 0.01: \nearrow (curved arrow pointing up and right)
- Under 0.03: \nearrow (curved arrow pointing up and right)
- Under 0.1: \nearrow (curved arrow pointing up and right)
- Under 0.3: \nwarrow (curved arrow pointing up and left)
- Under 1: \nwarrow (curved arrow pointing up and left)

Additional handwritten notes:

- Under 0.001: $3\times$
- Under 0.003: $\sim 3\times$
- Under 0.01: $3\times$
- Under 0.03: $\sim 3\times$
- Under 0.1: $\sim 3\times$



Machine Learning

Linear Regression with multiple variables

Features and
polynomial regression

Housing prices prediction

$$h_{\theta}(x) = \theta_0 + \theta_1 \times \boxed{\text{frontage}} + \theta_2 \times \boxed{\text{depth}}$$

x_1

x_2



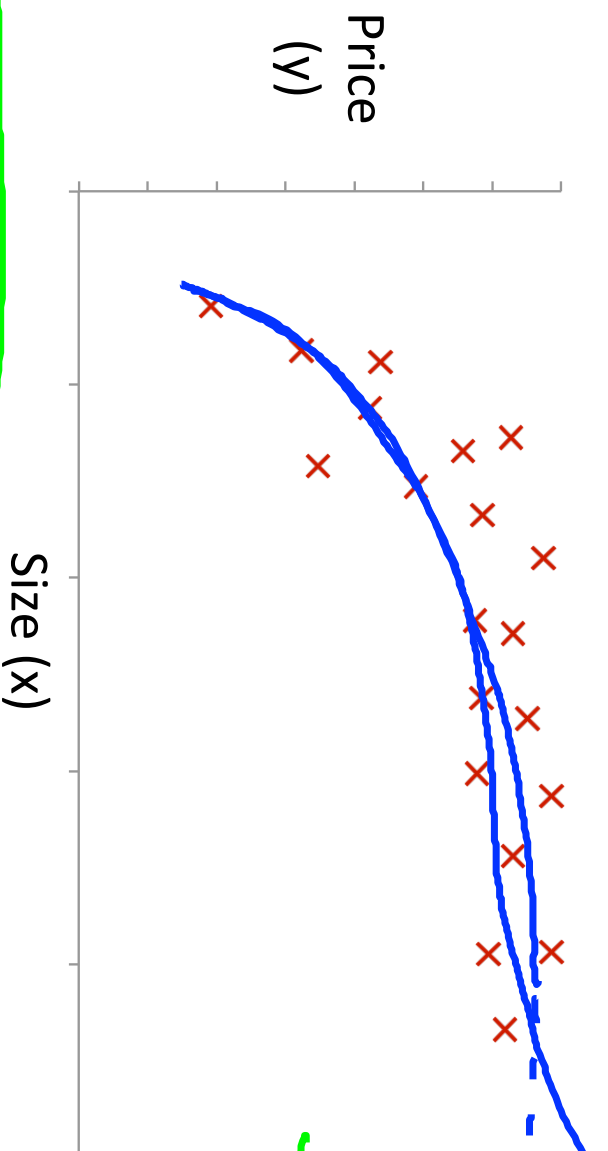
Area

$x = \text{frontage} \times \text{depth}$

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

\hookrightarrow land area

Polynomial regression



$$\rightarrow \theta_0 + \theta_1 x + \theta_2 x^2$$

$$\rightarrow \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3$$

Size (x)

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3$$

$$= \theta_0 + \theta_1(\text{size}) + \theta_2(\text{size})^2 + \theta_3(\text{size})^3$$

$$\rightarrow x_1 = (\text{size})$$

$$\rightarrow x_2 = (\text{size})^2$$


$$\rightarrow x_3 = (\text{size})^3$$

Size: 1-1000

Size²: 1-1000,000


Size³: 1-10⁹


Choice of features



→ $h_{\theta}(x) = \theta_0 + \theta_1(\text{size}) + \theta_2(\text{size})^2$

→ $h_{\theta}(x) = \theta_0 + \theta_1(\text{size}) + \theta_2\sqrt{(\text{size})}$






Machine Learning

Linear Regression with multiple variables

Normal equation

Gradient Descent




Normal equation: Method to solve for θ analytically.

Intuition: If 1D ($\theta \in \mathbb{R}$)

→ $J(\theta) = a\theta^2 + b\theta + c$

$\frac{\partial}{\partial \theta} J(\theta) = \dots \stackrel{\text{set}}{=} 0$

Solve for θ



$$\theta \in \mathbb{R}^{n+1} \quad J(\theta_0, \theta_1, \dots, \theta_m) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$\frac{\partial}{\partial \theta_j} J(\theta) = \dots \stackrel{\text{set}}{=} 0$ (for every j)

Solve for $\theta_0, \theta_1, \dots, \theta_n$

Examples: $m = 4$.

	Size (feet ²)	Number of bedrooms	Number of floors	Age of home (years)	Price (\$1000)
x_0	x_1	x_2	x_3	x_4	y
1	2104	5	1	45	460
1	1416	3	2	40	232
1	1534	3	2	30	315
1	852	2	1	36	178

$$\underline{X} = \begin{bmatrix} 1 & 2104 & 5 & 1 & 45 \\ 1 & 1416 & 3 & 2 & 40 \\ 1 & 1534 & 3 & 2 & 30 \\ 1 & 852 & 2 & 1 & 36 \end{bmatrix}$$

$m \times (n+1)$

$$\underline{y} = \begin{bmatrix} 460 \\ 232 \\ 315 \\ 178 \end{bmatrix}$$

m -dimensional vector

$$\theta = (X^T X)^{-1} X^T y$$

m examples $(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})$; n features.

$$\underline{x^{(i)}} = \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ x_2^{(i)} \\ \vdots \\ x_n^{(i)} \end{bmatrix} \in \mathbb{R}^{n+1}$$

$$X \begin{matrix} \text{(design matrix)} \\ m \times n \end{matrix} = \begin{bmatrix} \text{---} & (x^{(1)})^T & \text{---} \\ \text{---} & (x^{(2)})^T & \text{---} \\ \vdots & \vdots & \vdots \\ \text{---} & (x^{(m)})^T & \text{---} \end{bmatrix}$$

E.g. If $\underline{x^{(i)}} =$

$$\begin{bmatrix} 1 \\ x_1^{(i)} \end{bmatrix} \quad \underline{X} =$$

$$\begin{bmatrix} x^{(1)} & \text{---} \\ \vdots & \vdots \\ x^{(m)} & \text{---} \end{bmatrix}_{m \times (n+1)} \quad \underline{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix}_{m \times 1}$$

$$\theta = \frac{1}{(X^T X)^{-1} X^T y}$$

←

$(X^T X)^{-1}$ is inverse of matrix $X^T X$.

Set A : $X^T X$

$$(X^T X)^{-1} = A^{-1}$$

Octave: `pinv(X' * X) * X' * y`

$$\text{pinv}(X^T * X) * X^T * y$$

$$\theta = \frac{1}{(X^T X)^{-1} X^T y} \min J(\theta)$$

X' X^T

~~Feature Scaling~~

$$0 \leq x_1 \leq 1$$

$$0 \leq x_2 \leq 1000$$

$$0 \leq x_3 \leq 10^{-5}$$

✓

m training examples, n features.

Gradient Descent

- • Need to choose α .
- • Needs many iterations.
- Works well even when n is large.

→ $n \approx 10^6$


← -

Normal Equation

- • No need to choose α .
- • Don't need to iterate.
- Need to compute $(\boxed{X^T X})^{-1}$ $n \times n$ $O(n^3)$
- • Slow if n is very large.

$n \approx 100$
 $n \approx 1000$

- - - $n \approx 10000$



Machine Learning

Linear Regression with multiple variables

Normal equation and non-invertibility (optional)

Normal equation

$$\theta = (X^T X)^{-1} X^T y$$

$$\underline{X^T X}$$

- What if $X^T X$ is non-invertible? (singular / degenerate)

- Octave: `pinv (X' * X) * X' * y`

①

pinv
→ inv

What if $X^T X$ is non-invertible?

- Redundant features (linearly dependent).

E.g. $x_1 = \text{size in feet}^2$

$$1 \text{ m} = 3.28 \text{ feet}$$

~~$x_2 = \text{size in m}^2$~~

$$x_1 = (3.28)^2 x_2$$

$$\rightarrow n = 10 \leftarrow$$

$$\rightarrow n = 100 \leftarrow$$

$$\Theta \in \mathbb{R}^{101}$$

- Too many features (e.g. $m \leq n$).

- Delete some features, or use regularization.

\downarrow later